

Dear Hiring Team,

I'm applying for the Pricing/Revenue-ML lead role, and I want to be direct: I don't have 5+ years of experience or a pricing/revenue track record. What I can offer is rigorous ML engineering practice and a genuine comfort with the evaluation discipline this role is built around.

My thesis is fundamentally about getting evaluation right before trusting a system: I formalized ontology repair as an MDP, built a closed-loop pipeline validated against two independent OWL reasoners (Konclude, HermiT) rather than a single metric, and tracked repair success (88-96%) and step efficiency (5.1 steps/episode) across violation types - the kind of "don't just look good, perform reliably" standard this role asks for. I'm used to thinking about failure modes before I trust a number.

At Fraunhofer IGCV, I worked across the full ML lifecycle - data ingestion, feature engineering, model training, Docker-based deployment, and monitoring - for a live manufacturing system, and combined causal discovery (DAGMA) with graph neural networks specifically to avoid mistaking correlation for the real signal. That mindset - being suspicious of a model that "just" performs well - is the same one this role needs around leakage and bias.

My Python and SQL are solid and production-oriented (Fraunhofer, Accenture), though I haven't worked with dbt, Snowflake, or Metabase specifically, and I don't have pricing or revenue-management domain experience yet. My German is B1, not fluent.

I'm early in my career relative to what you're asking for, but the engineering discipline - reproducibility, honest evaluation, treating a model as unproven until tested - is already how I work. If that outweighs the years on paper for you, I'd welcome a conversation.

Best regards,

Resham Joshi